Efficient Calculation of Interaction Features on Sparse Matrices

Andrew Nystrom

Abstract

FILL THIS IN

1 Introduction

Introduction Interaction features are a way of capturing correlations between features in a machine learning setting. A feature vector $\vec{x}$ of dimensionality D has second degree interaction features $\{x_a \cdot x_b : a, b \in \{0, 1, ..., D-1\} \wedge a < b\}$, so a D dimensional vector has $\binom{D}{2} = \frac{D^2-D}{2}$ second degree interaction features. A naive approach to calculating these features is to simply iterate through the combinations of the column indices. For a sparse vector, many of the resulting interaction features would be zero, and could therefore be ignored. This work describes a method to efficiently calculate second degree interaction features for a sparse matrix that has time and space complexities that decrease quadratically with the density of the input matrix with respect to the naive approach.

2 Approach

Let the list of nonzero columns for a given row $\vec{x}$ be denoted by N_{zc}. The nonzero second degree interaction features are simply the products of all combinations of two elements whose columns are in N_{zc}. However, to properly place an interaction feature into the correct column, a mapping from the column index pairs of N_{zc} into the columns of the interaction matrix is

needed. The mapping is from the space (a, b) where a,b are in $1, 2, ..., D$ onto the space $1, 2, ..., \frac{D^2-D}{2}$. This is isomorphic to mapping the coordinates of the upper triangle of a matrix onto a flat list. The following is a proof by construction for such a mapping.

INSERT JOHN'S PROOF HERE

With this mapping, an algorithm for generating second degree interaction features on a matrix A can be formulated as follows:

SPARSE INTERACTION(A)
    $\text{map}(b, a) = \frac{2Da-a^2+2b-3a-2}{2}$
    B = Compressed Sparse Row Matrix of size $N \times \frac{D^2-D}{2}$
    for *row* in A
        N_{zc} = nonzero columns of *row*
        for each combination a, b of elements of N_{zc}
            $k = \text{map}(b, a)$
            i = index of *row*
            $B[i, k] = row[a] \cdot row[b]$

3 Complexity Analysis

Assume that A is a matrix with sparsity $0 < d < 1$, N rows, and D columns. Finding interaction features with the proposed algorithm has time and space complexity $\text{O}(dND^2)$, whereas a naive approach of using non-sparse matrices and multiplying all column combinations has time and space complexity $\text{O}(ND^2)$. The algorithm is therefore an improvement by a factor of the density factor of A.

This can represent a large gain in speed and time. For example, the 20 Newsgroups dataset has density d of 0.12 when its unigrams are represented in a vector space model. This means the proposed approach would take less than $\frac{1}{8}$ time and memory.

The real benefit of this method is revealed when the average complexity is analysed. The number of interactions calculated for a given row are $\binom{|N_{zc}|}{2}$. If the matrix has density d, then on average, $N_{zc} = Dd$, so the number of interaction features calculated in total is

$$N\binom{dD}{2} = N\frac{(Dd)!}{2!(Dd-2)!}$$

$$= N\frac{(D^2d^2 - Dd)}{2}$$

This means that the average complexity decreases quadratically with the density.

4 Future Work

The approach for generating second degree interaction features required a mapping from combinations of two to the space $1, 2, ..., \frac{D^2-D}{2}$, which is isomorphic to a mapping from the indices of an upper triangular matrix to the indices of a flat list of the same size. To generate third degree interaction features, a mapping from combinations of three (a, b, c) to the space $1, 2, ...\frac{D^3-3D^2+2D}{6}$ (which is $\binom{D}{3}$), or the upper 3-simplex of a tensor to a flat list of the same size $\frac{D^3-3D^2+2D}{6}$ would be required. In general, for interaction features of degree k, the upper k-simplex of a k-dimensional tensor must be mapped to the space $1, 2, ...\frac{D!}{k!(D-k)!}$. A similar approach for finding these mappings could be taken as the one used here for $k = 2$.

Motivation for deriving mapping functions for higher orders of interaction features is that the average complexity of generating degree k interaction features is $N\binom{Dd}{k}$, which decreases polynomially with respect to k compared to generating the features naively.